


# Kubernetes

*Basics*






Everyone's container journey starts with one container....



At first the growth is easy to handle....





But soon you have many applications, many instances...

And that is why there is container orchestration



# What is container orchestration?

Management of the deployment, placement, and lifecycle of workload containers

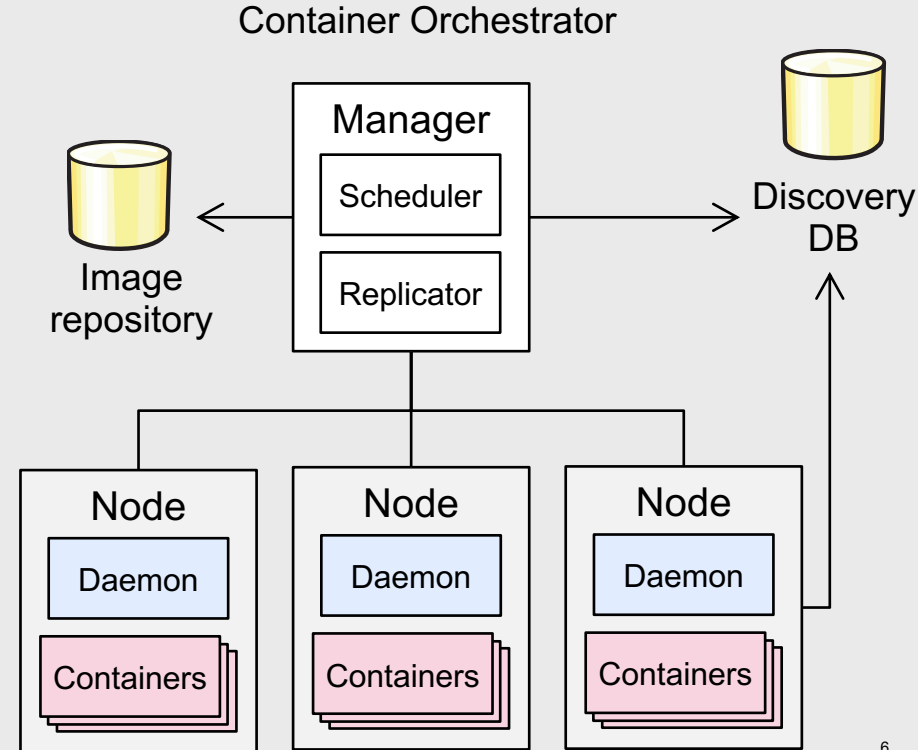
Cluster management creates unified targets for varied workload

Scheduling intelligently distributes containers across nodes

Service discovery knows where containers are located and provides a method to direct requests to them

Replication allows the workload to be scaled

Health management creates a method to ensure the application is assured to be viable by allowing unhealthy containers to be replaced



# Kubernetes and Container Orchestration

## Kubernetes is a Container Orchestrator

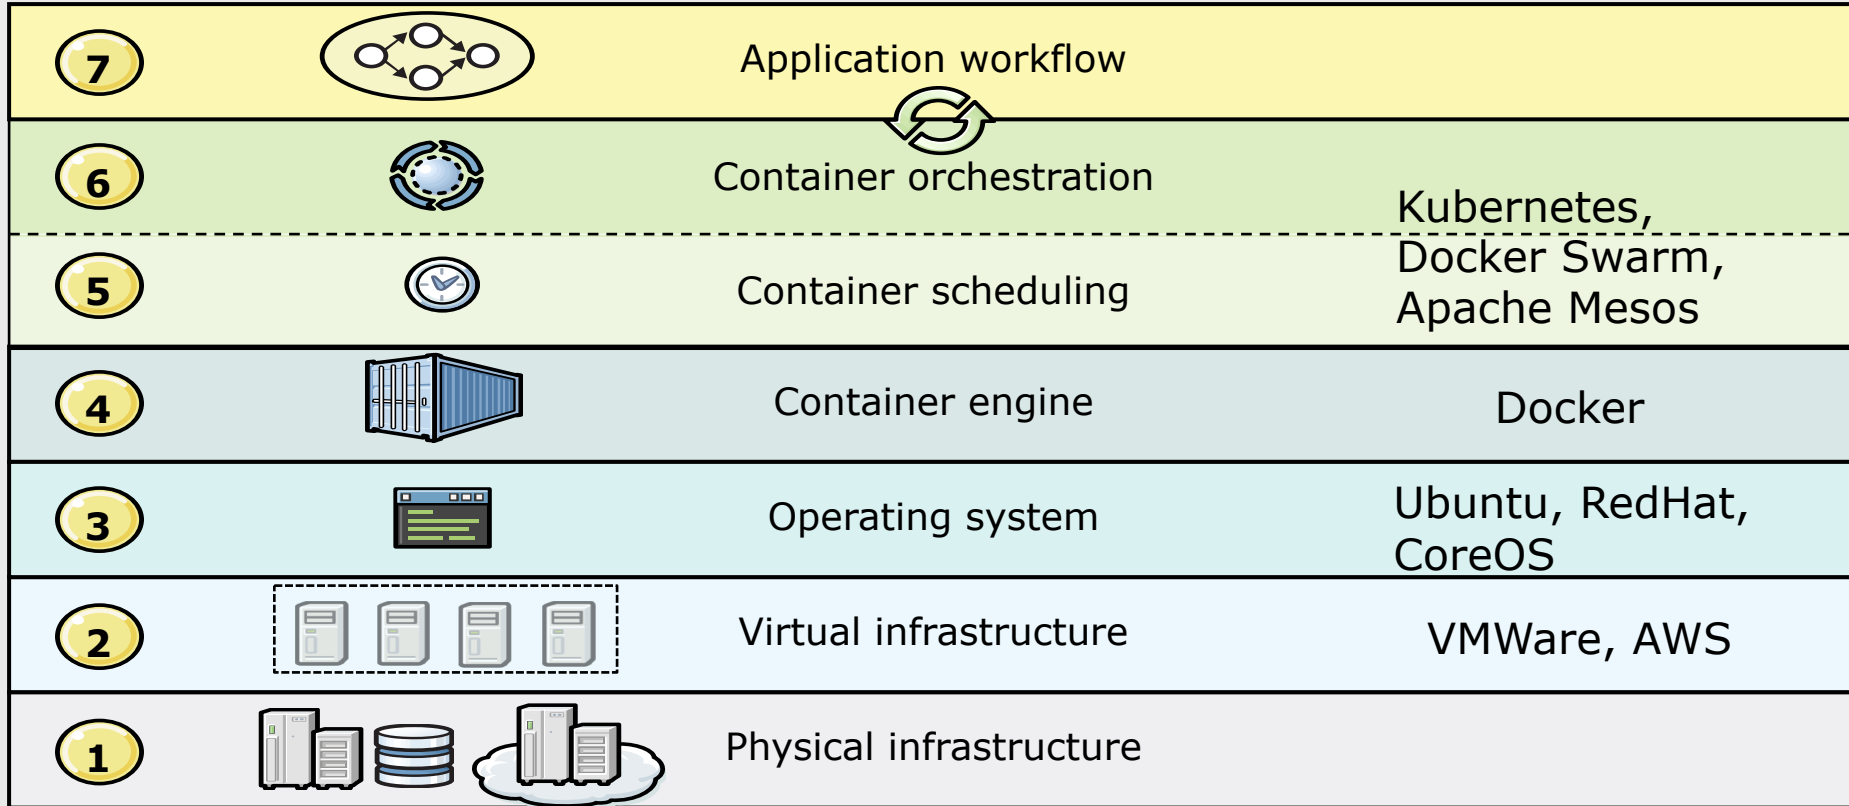
- Provisions, manages, and scales containerized applications
- Supports:
  - Automated scheduling and scaling
  - Zero downtime deployments
  - High availability and fault tolerance
  - A/B deployments
- Manage infrastructure resources needed by applications
  - Volumes
  - Networks
  - Secrets
  - And many many many more..
- Declarative model
  - Provide the "desired state" and Kubernetes will make it happen



## Kubernetes fun facts:

- Builds upon 15 years of experience of running production workloads at Google
- Open Governance via CNCF
- Adopted by IBM, Amazon, Microsoft, Red Hat, Google,
- Means “helmsman” in Greek

# Container ecosystem






# Kubernetes

*The Building Blocks*

# Immutability

Build Once - Deploy Everywhere



The same container image is built once and is moved between environments



Dev

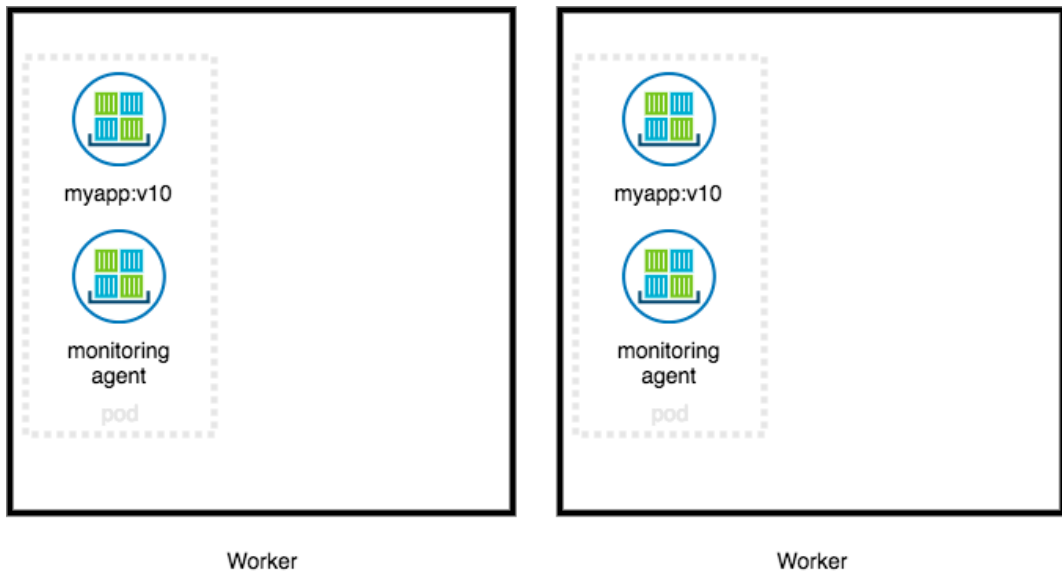


Prod

# Pod

A single unit of work in Kubernetes, which may consist of one or more containers

All containers in a pod are co-located and co-scheduled, and share the kernel namespace (process, storage, network, etc.)



# Pod Health Checking

Pods are automatically kept alive by “process check” checking the basic status of the main process for the application

To go beyond this Kubernetes allows you to create a liveness probe to provide additional means for identifying health.

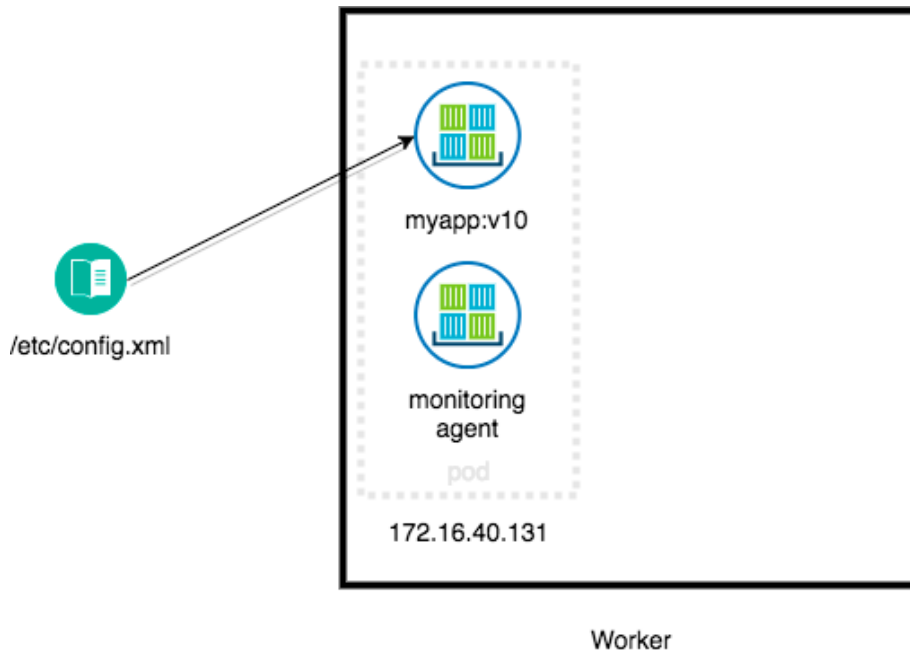
```
apiVersion: v1
kind: Pod
metadata:
  name: pod-with-http-healthcheck
spec:
  containers:
    - name: nginx
      image: nginx
      # defines the health checking
      livenessProbe:
        # an http probe
        httpGet:
          path: /_status/healthz
          port: 80
        # length of time to wait for a pod to initialize
        # after pod startup, before applying health checking
        initialDelaySeconds: 30
        timeoutSeconds: 1
      ports:
        - containerPort: 80
```

# Config Maps & Secrets

Share and store configurations, credentials and more

Store the configurations and secrets (credentials, certificates) in the K8s environment and mount them to the local filesystem within container(s)

The container image can move un-changed between environments (i.e. container immutability)

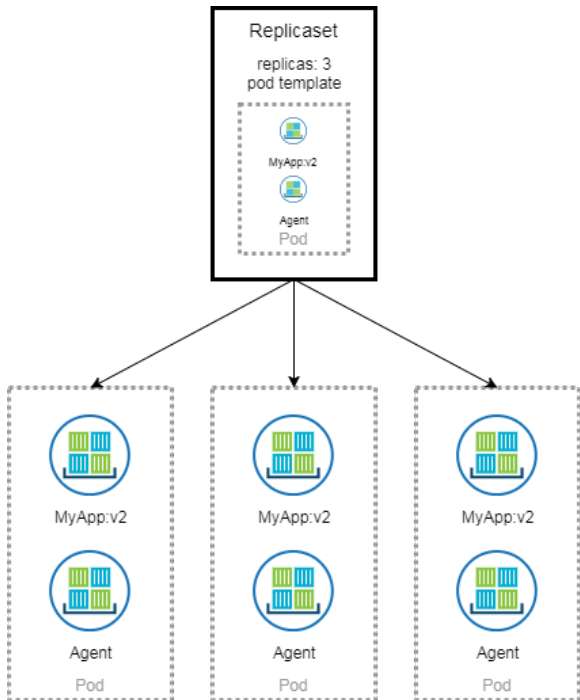


# Replicaset

Scale pods horizontally and provide resiliency

Replicasets run one-to-many instances of the desired pod

When possible the replica pod should be stateless or near-stateless




```
apiVersion: apps/v1beta2 # for versions before 1.8.0 use apps/v1beta1
kind: ReplicaSet
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  # this replicas value is default
  # modify it according to your case
  replicas: 3
  selector:
    matchLabels:
      tier: frontend
    matchExpressions:
      - {key: tier, operator: In, values: [frontend]}
  template:
    metadata:
      labels:
        app: guestbook
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google_samples/gb-frontend:v3
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          env:
            - name: GET_HOSTS_FROM
              value: dns
              # If your cluster config does not include a dns service, then to
              # instead access environment variables to find service host
              # info, comment out the 'value: dns' line above, and uncomment the
              # line below.
              # value: env
      ports:
        - containerPort: 80
```

# Deployments

Deployments manage rolling updates to ReplicaSets and StatefulSets

When a new version of the application is available, the Deployment provides the ability to scale down the previous version of the application and scale up the new version in a controlled fashion with zero downtime

Enables rollback in the case of failure




# Service Discovery

Kubernetes has an internal DNS that is used as a Service Registry.

A Service resource in Kubernetes results in an entry in the internal DNS

By default, a Service points to an internal Cluster IP that load balances between a set of healthy running pods

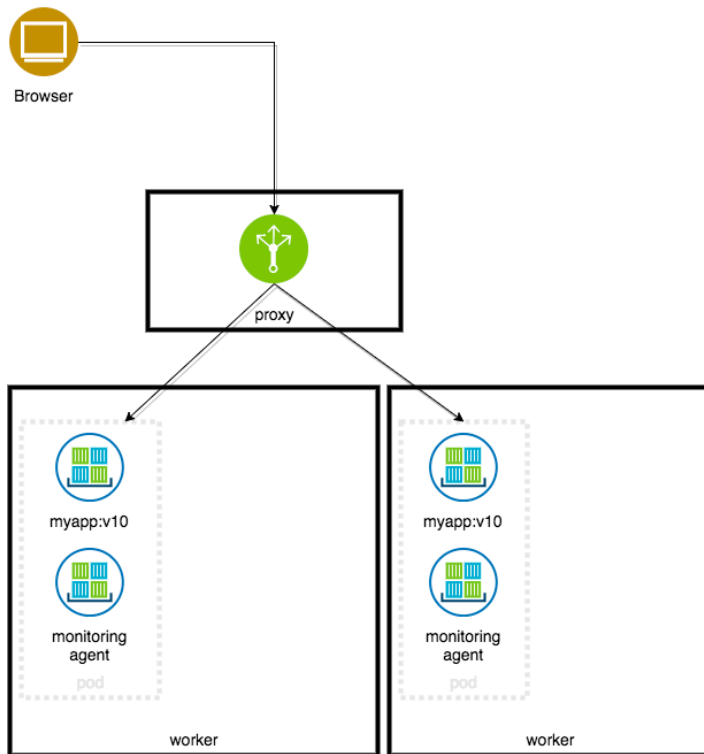




# Ingress Resources

External access to applications running in Kubernetes may be enabled through Ingress resources and proxy nodes

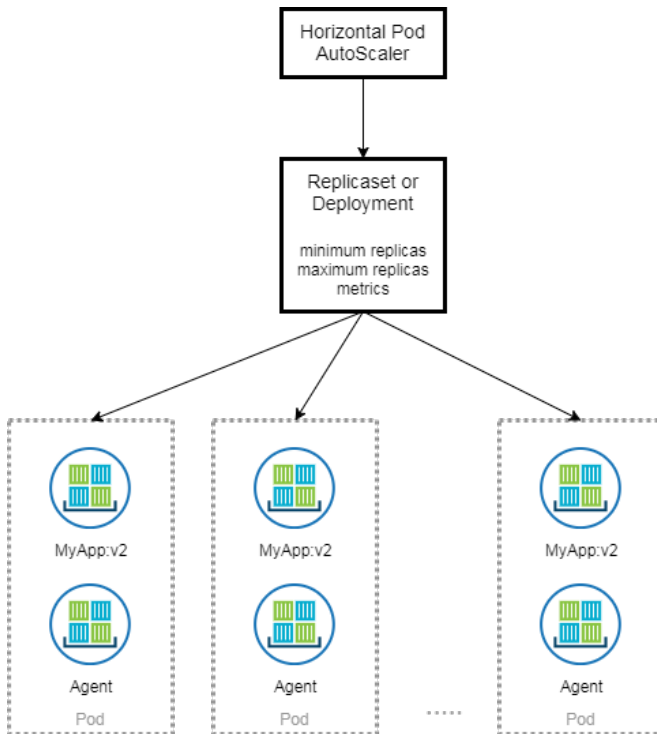
The proxy node(s) in ICP  
expose services defined in  
Kubernetes



# More on Scaling

## Horizontal Pod Auto-scaling (HPA)

Allows you to scale the number of running pods in a replicaset based upon resource (or application custom) metrics



```
apiVersion: autoscaling/v2beta1
kind: HorizontalPodAutoscaler
metadata:
  name: php-apache
  namespace: default
spec:
  scaleTargetRef:
    apiVersion: apps/v1beta1
    kind: Deployment
    name: php-apache
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      targetAverageUtilization: 50
status:
  observedGeneration: 1
  lastScaleTime: <some-time>
  currentReplicas: 1
  desiredReplicas: 1
  currentMetrics:
  - type: Resource
    resource:
      name: cpu
      currentAverageUtilization: 0
      currentAverageValue: 0
```

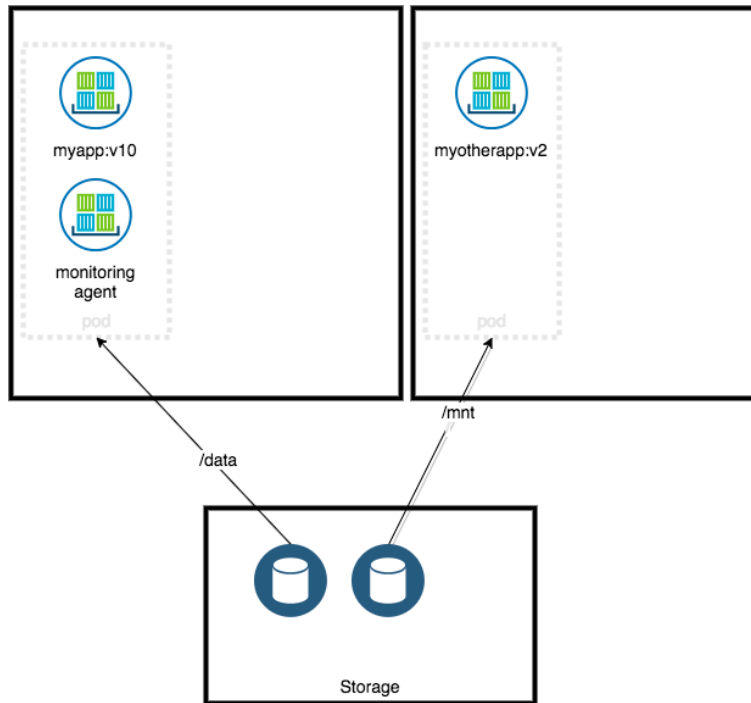
# Persistence & Storage

There are many types of persistent storage and many provider options

Some pods must be able to persist data so that if Kubernetes restarts them on the same or another node data loss is avoided

Kubernetes will re-attach the shared storage when the pod (re)starts

Storage providers support different retention and recycling policies and the definitions of these are not universal



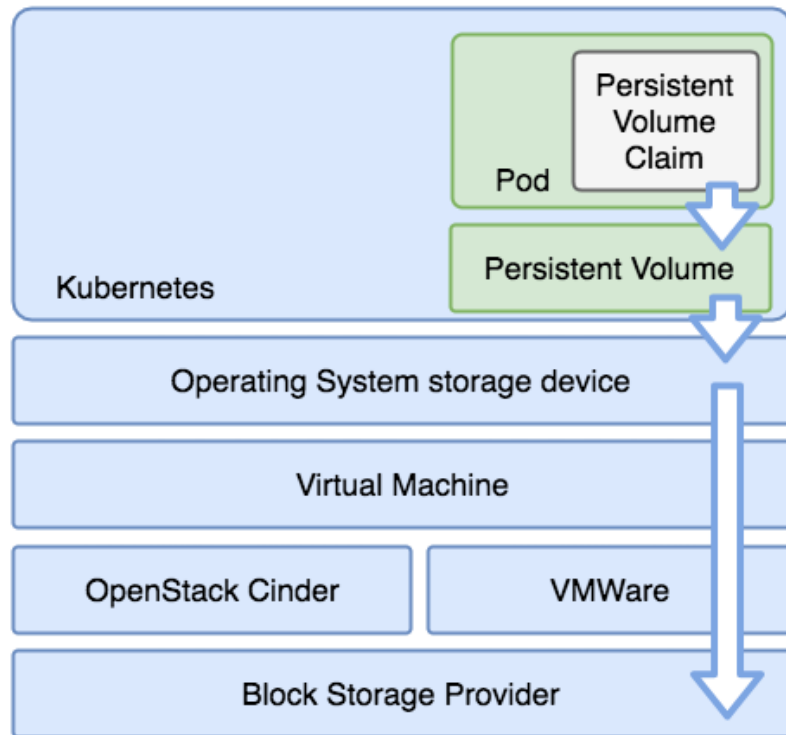
# Persistent Storage

## Solution components

**Persistent Volume** is a storage resource within the cluster. PVs have a lifecycle independent of any individual pod that uses it. This API object encapsulates the details of the storage implementation or cloud-provider-specific storage system.

A **Persistent Volume Claim** is a storage request, or claim, made by the developer. Claims request specific sizes of storage, as well as other aspects such as access modes.

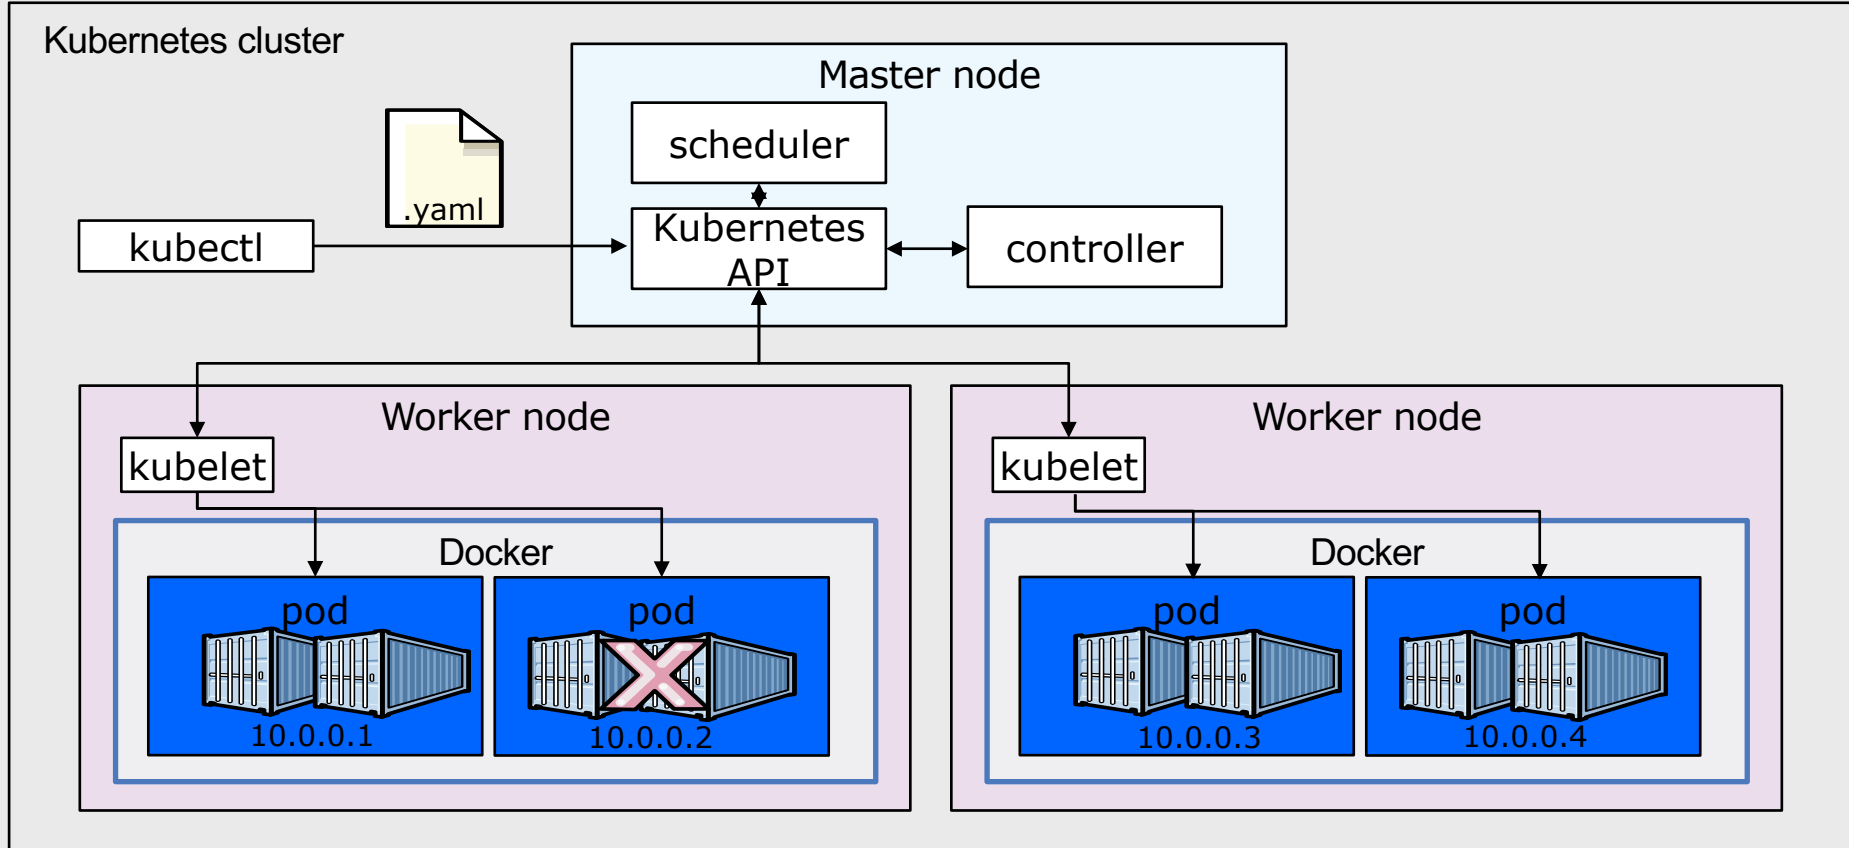
A **StorageClass** describes an offering of storage and allow for the dynamically provisioning of PVs and PVCs based upon these controlled definitions.



# Kubernetes

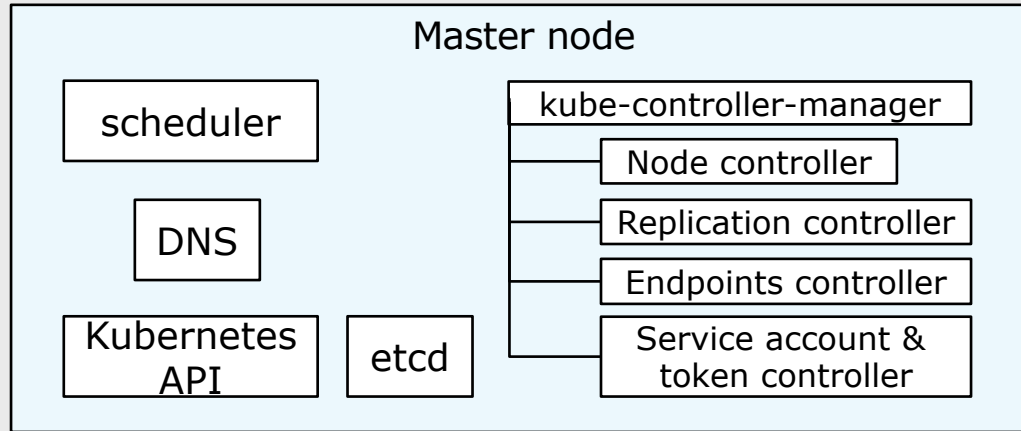
*The Cluster*

# Kubernetes cluster architecture



# Master Node components

- Make scheduling decisions for the cluster, and respond to cluster events, like a node failure
- Can run on any node in the cluster, but typically all master components run on the same virtual machine (vm), and do not run any container apps on that vm



# Master Node Components

## Etcd

- A highly-available key value store
- Stores all cluster data

## API Server

- Exposes API for managing Kubernetes
- Used by kubectl CLI

## Scheduler

- Selects the worker node for each pods runs


## Controller manager

- Daemon that runs controllers (background threads that handle routine tasks in the cluster)
- Node Controller – Responsible for noticing and responding when nodes go down
- Endpoints Controller – Populates the Endpoints object (joins services and pods)
- Service Account and Token Controllers – Create default accounts and API access tokens for new namespaces



# Worker Node Components

- Provide the Kubernetes runtime environment; run on every node
- Maintain running pods



# Kubernetes Clients (CLI and Dashboard)

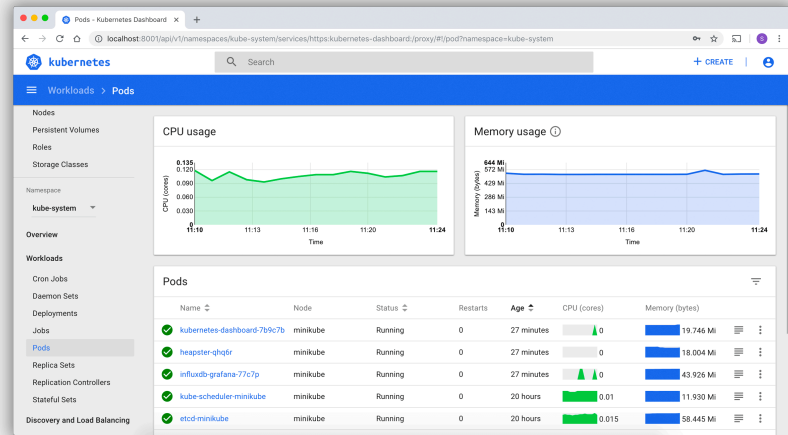
## Kubernetes CLI

- Directly manipulate YAML
  - `kubectl (create|get|apply|delete) -f myResource.yaml`
- <https://kubernetes.io/docs/tasks/tools/install-kubectl>



## Kubernetes Dashboard

- Another way to view and modify resources



# Kubernetes in Action!

1. User via "**kubectl**" deploys a new application
2. API server receives the request and stores it in the DB (etcd)
3. Watchers/controllers detect the resource changes and act upon it
4. ReplicaSet watcher/controller detects the new app and creates new pods to match the desired # of instances
5. Scheduler assigns new pods to a kubelet
6. Kubelet detects pods and deploys them via the container runtime (e.g. Docker)
7. Kubeproxy manages network traffic for the pods – including service discovery and load-balancing

